

Impact Statement — African Folktales SLM

Our model touches something older and more sensitive than most NLP applications: living oral traditions that belong to specific communities, not to the internet at large. That changes what "impact" means here. The applications are fairly concrete — a system that can retrieve, contextualize, and translate folktales for classrooms, diaspora families, and cultural researchers. The implications, though, run in two directions at once, and we don't think it's honest to pretend otherwise.

On the positive side, a well-built version of this tool could genuinely help. A Yoruba child growing up in London or a Zulu teenager living far from home could access stories from their own lineage in a form that's searchable, translated, and tied back to a named community rather than a flattened "African folklore" bucket. Done carefully, this strengthens intergenerational transmission at a moment when fewer children are hearing these tales directly from elders.

But the same properties that make the system useful also make it risky. Any model trained on oral material is at risk of doing exactly what it's meant to prevent: flattening distinct traditions into a generic pan-African "trickster tale" voice, stripping stories of the regional texture that gives them meaning. There's also a real possibility that easy digital access reduces incentive to seek out actual storykeepers — the tool could quietly substitute for the relationship it's supposed to support. Misattribution is another concrete risk: a story labeled to the wrong ethnolinguistic group isn't a minor error, it's a form of erasure, especially for smaller groups already at risk of being overshadowed by dominant languages like Swahili or Hausa.

We're also not fully certain our consent-first pipeline scales. Opt-in consent from a handful of early sources is easy to secure; verifying it across dozens of oral sources, especially from communities with limited digital infrastructure, is harder, and we don't have a complete answer yet for how attribution holds up once a story has been paraphrased or translated by the model.

Going forward, we think the most useful next steps aren't purely technical: partnering with linguists to build per-language accuracy audits, establishing a standing review process rather than a one-time approval, and commissioning independent research on whether tools like this actually increase engagement with living storytellers or quietly replace them. We'd rather surface that last question honestly than assume the answer.

